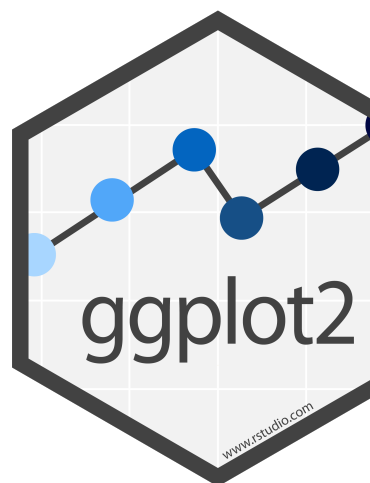


Data Visuali

JSC 370: Data Sc

January 29, 20

Background



This lecture provides an introduction to `ggplot2`, which provides vastly better graphics options than R's default

Background

`ggplot2` is part of the Tidyverse. The tidyverse is a collection of R packages designed for data science. All packages follow the same design philosophy, grammar, and data structure. Visit www.tidyverse.org/

Core tidyverse includes `dplyr`, which provides data manipulation.

It also includes `stringr` which we will use in a moment. We saw in lab that helps with dates and times.

```
library(tidyverse)
library(nycflights13)
```

```
library(kableExtra)
```

Layers, dplyr and pipes

- We should take a step back and discuss
- `ggplot2` behaves very similarly to `dplyr`. The
- The first argument in `dplyr` is always a data frame (or a `data.table`).
- Subsequent arguments can also be thought of as actions to be taken on the data (verbs).
- The output is always a new data frame.
- Layers are connected by a pipe, which un

- The new pipe is `|>`, which works similarly likely only noticeable by expert users.

The Pipe `%>%` and now `|>`

- The pipe passes the object on its left hand of the function on the right hand side.
- We can kind of think of it as saying 'then'

Flights data

To illustrate many of today's examples we will use the `nycflights13` library. They are all flights that departed from New York City (JFK, LGA, EWR) in 2013.

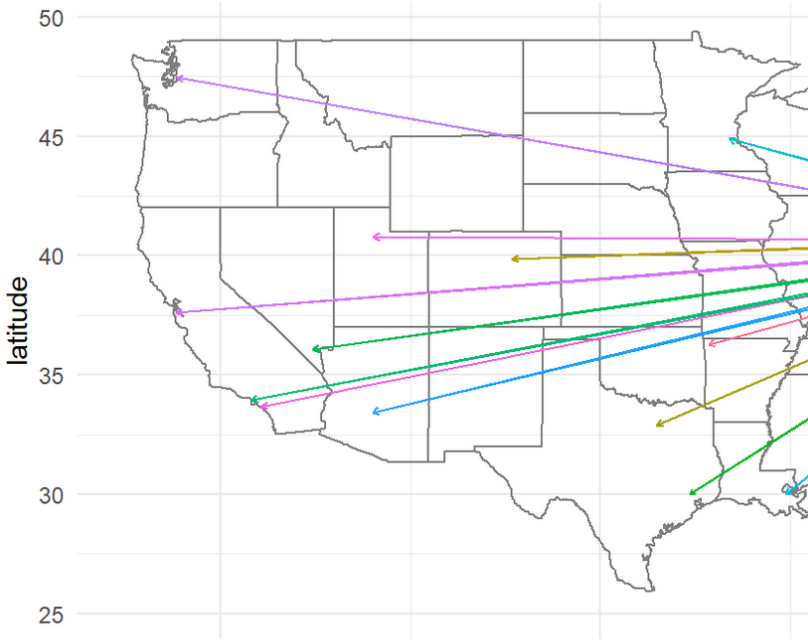
```
names(flights)
```

```
## [1] "year"          "month"         "day"
## [5] "sched_dep_time" "dep_delay"     "arr_time"
## [9] "arr_delay"     "carrier"       "flight"
## [13] "origin"        "dest"         "air_time"
## [17] "hour"         "minute"       "time_hour"
```

```
dim(flights)
```

[1] 336776 19

Flights data



-120

-100

longitude

The pipe %>%

- An example using pipes: subset the flight data, calculate the mean arrival delay times by year, month, carrier.
- We need to start with the data, filter to get the data we want, group_by to prepare the groups that we want to summarize to take the mean (or whatever function we are interested in).


```
nycflights13::flights %>%  
  filter(dest == "LAX") %>%  
  group_by(year, month, day) %>%  
  summarize(  
    arr_delay = mean(arr_delay, na.rm = TRUE)  
  )
```

The new pipe |>

Let's do the same thing with the new pipe:

```
nycflights13::flights |>
  filter(dest == "LAX") |>
  group_by(year, month, day) |>
  summarize(
    arr_delay = mean(arr_delay, na.rm = TRUE)
  )
```

A few coding style tips

- Variable names (those created by `<-` and `summarize()`) should use only lowercase
- Use `_` to separate words within a name
- `%>%` or `|>` should always have a space before followed by a new line
- After the first step, each line should be indented
- Note on the previous slides that had long function may not all fit on one line, put each line on a new line and indent.

```
short_flights <- flights |>
  filter(air_time < 60)
```

Flights data in more detail

The `nycflights13` library also provides hourly data for three NYC airports (the origin). Let's join the `weather` data so we can look at more interesting relationships.

```
names(weather)
```

```
## [1] "origin"      "year"        "month"       "day"
## [6] "temp"        "dewp"        "humid"       "wind_dir"
## [11] "wind_gust"   "precip"      "pressure"    "visib"
```

```
dim(weather)
```

```
## [1] 26115      15
```

Flights data in more detail

Looks like we can join these datasets on year, month, day, and hour (which is the origin airport). We can examine flight delays and weather at the origin airport.

A `left_join` will keep all of the observations in `x` (flights) with the observations in `y` (weather at origin). We can examine flight delays originating at the 3 NYC airports on a given day. We want to keep the resolution of the `x` dataset.

```
flights_weather <-  
  left_join(  
    flights, weather, by = c("year", "month", "day", "hour")  
  )
```

Flights data in more detail

```
head(flights_weather)
```

```
## # A tibble: 6 × 29
##   year month   day dep_time sched_dep_time dep_delay a
##   <int> <int> <int>   <int>         <int>      <dbl>
## 1  2013     1     1     517           515         2
## 2  2013     1     1     533           529         4
## 3  2013     1     1     542           540         2
## 4  2013     1     1     544           545        -1
## 5  2013     1     1     554           600        -6
## 6  2013     1     1     554           558        -4
## # i 21 more variables: arr_delay <dbl>, carrier <chr>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <
## #   hour <dbl>, minute <dbl>, time_hour.x <dtm>, temp
## #   humid <dbl>, wind_dir <dbl>, wind_speed <dbl>, wind
## #   precip <dbl>, pressure <dbl>, visib <dbl>, time_hou
```

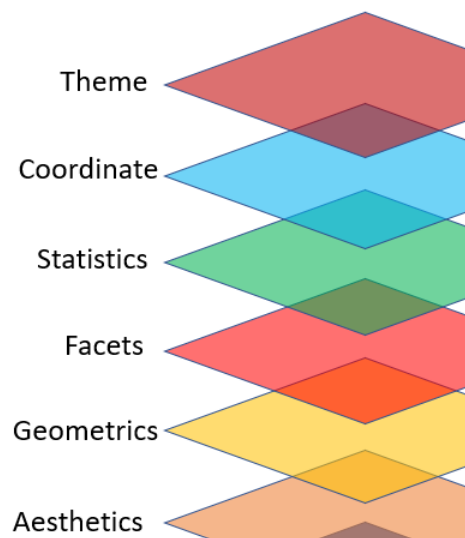
Daily flights data

Let's make a more manageable sized dataset
month, day, and origin airport

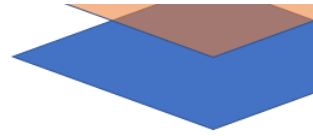
```
flights_weather_day <-  
  flights_weather |>  
  group_by(year, month, day, origin) |>  
  summarize_at(  
    vars(dep_delay, arr_delay, temp, dewp, humid, wind_dir,  
          wind_speed, wind_gust, precip, pressure, vis)  
  )
```

Visualizations

`ggplot2` is designed on the principle of adding



Data



Layers in ggplot2

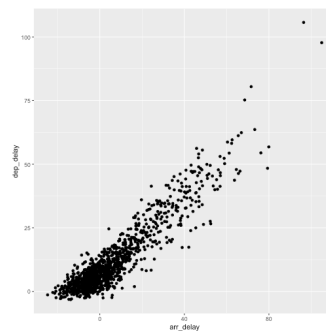
- With `ggplot2` a plot is initiated with the function `ggplot()`
- The first argument of `ggplot()` is the data
- We add aesthetics is always paired with `aes()`
- The `aes()` mapping takes the x and y axes
- Layers are added to `ggplot()` with `+`
- Layers include `geom` functions such as `point`

```
ggplot(data = data, mapping = aes(mappings)) +  
  geom_function()
```

Basic scatterplot

The first argument of `ggplot()` is the dataset to plot. With the `+` you add one or more layers.

```
ggplot(data = flights_weather_day, mapping = aes(x = arr_delay, y = dep_delay)) +  
  geom_point()
```

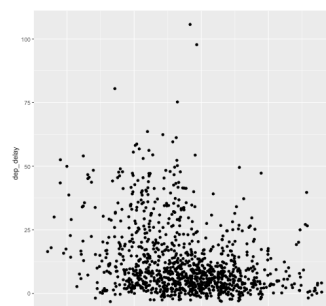


As expected, we see that if a flight has a late c

Another basic scatterplot

We can drop `data =` and `mapping =`. Now let's s
between departure delays and pressure (low
precipitation, high pressure means better we

```
ggplot(flights_weather_day, aes(x = pressure, y = dep_del  
  geom_point()
```





Adding to a basic scatterplot

- `geom_point()` adds a layer of points to your plot.
- `ggplot2` comes with many geom functions to add different type of layer to a plot.
- Each geom function in `ggplot2` takes a mapping argument.
- This defines how variables in your dataset are mapped to the properties of the geom.
- The mapping argument is always paired with the `aes()` arguments of `aes()` specify which variables are mapped to the properties.
- One common problem when creating a ggplot is that the

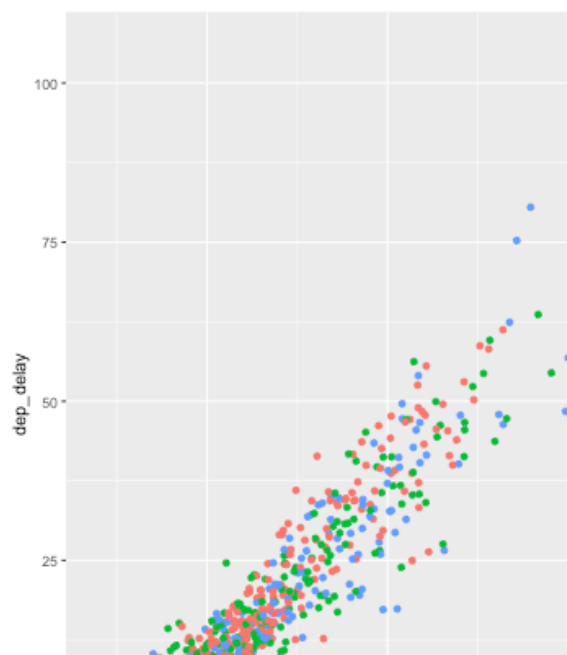
in the wrong place: it has to come at the end

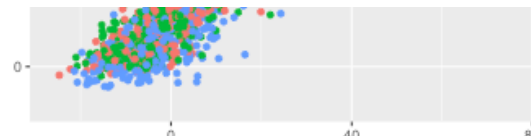
Coloring by a variable - using

You can convey information about your data in your plot to the variables in your dataset. For example, you can color the points of your plot by the class variable to represent the class. `geom_point()` chooses colors, and adds a legend, automatically.

```
ggplot(flights_weather_day, aes(x = arr_delay, y = dep_delay)) +  
  geom_point()
```

Coloring by a variable - using



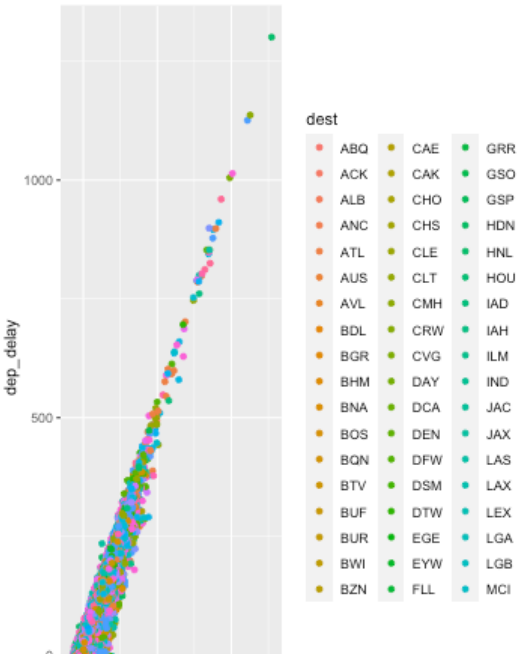


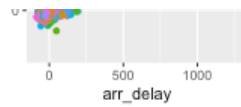
Coloring by a variable - using

Note when there are a lot of classes or groups distinguished well

```
ggplot(flights_weather, aes(x = arr_delay, y = dep_delay,  
  geom_point()
```

Coloring by a variable - using



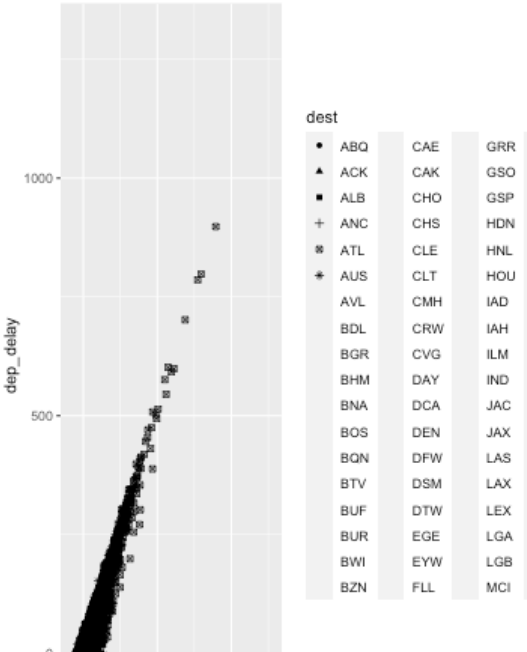


Determining point type using

By default ggplot uses up to 6 shapes. If there is not plotted!! (At least it warns you.)

```
ggplot(flights_weather, aes(x = arr_delay, y = dep_delay,  
  geom_point()
```

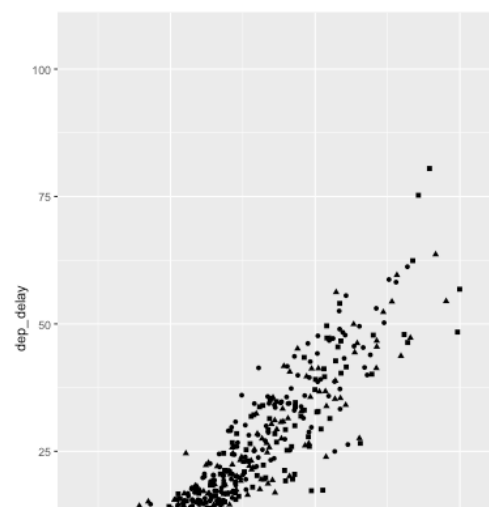
Determining point type using

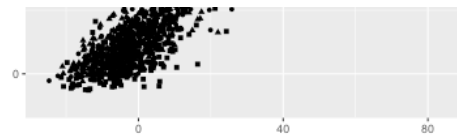




Determining point type using

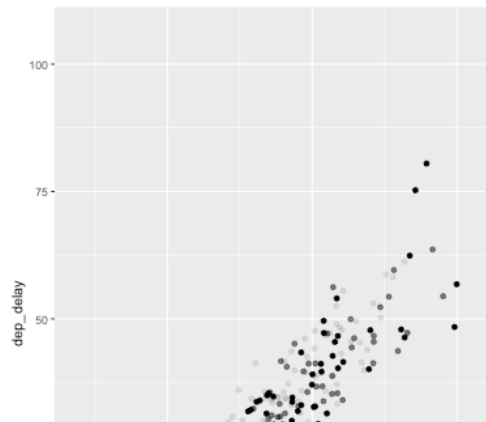
```
ggplot(flights_weather_day, aes(x = arr_delay, y = dep_delay))  
  geom_point()
```

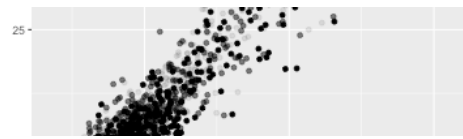




Controlling point transparency "alpha" aesthetic

```
ggplot(flights_weather_day, aes(x = arr_delay, y = dep_delay))  
  geom_point(alpha = 0.5)
```

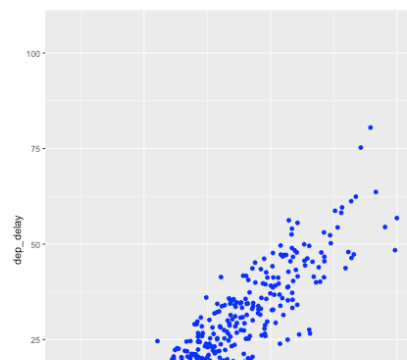


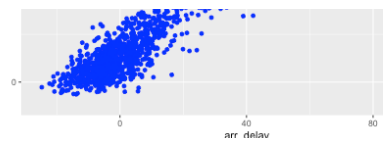


Manual control of aesthetics

To control aesthetics manually, we do it in `geom_point()` name outside `aes()`

```
ggplot(flights_weather_day, aes(x = arr_delay, y = dep_delay)) +  
  geom_point(color = "blue")
```





Summary of aesthetics

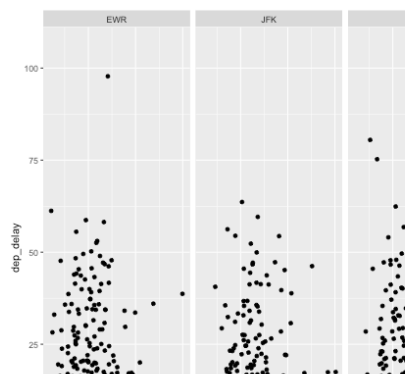
The various aesthetics...

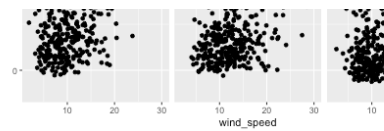
Code	Description
x	position on x-axis
y	position on y-axis
shape	shape
color	color of elements
fill	color inside elements
size	size
alpha	transparency
linetype	type of line

Facets 1

Facets are particularly useful for categorical v

```
ggplot(flights_weather_day, aes(x = wind_speed, y = dep_delay)) +  
  geom_point() +  
  facet_wrap(~origin, nrow = 1)
```

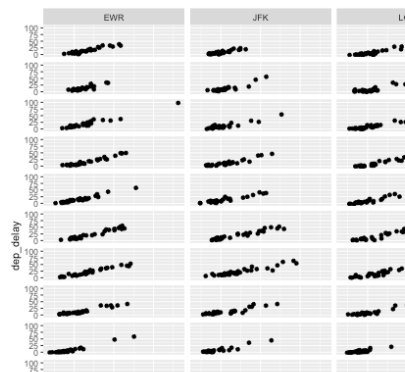


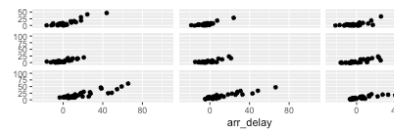


Facets 2

Or you can facet on two variables...

```
ggplot(flights_weather_day, aes(x = arr_delay, y = dep_delay)) +
  geom_point() +
  facet_grid(month ~ origin)
```



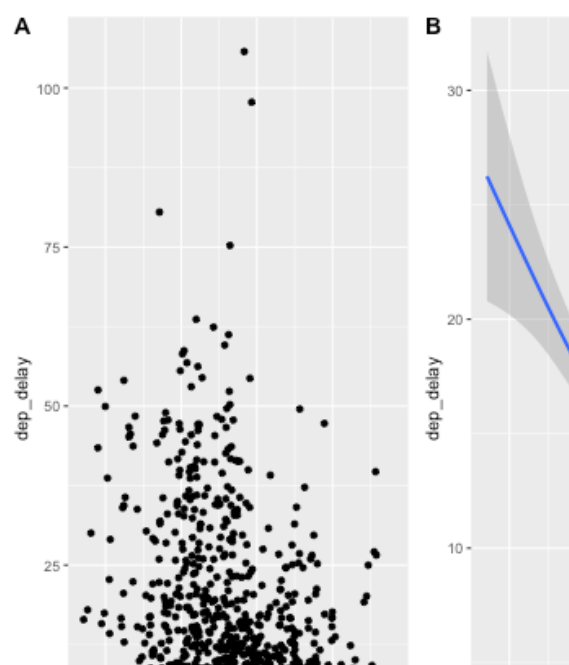


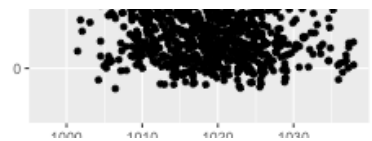
Geometric Objects 1

Geometric objects are used to control the type of plot (e.g., scatterplot, lineplot, facet, etc.) when plotting a smoothed line fitted to the data (arr_delay vs. pressure, side plots).

```
library(cowplot)
scatterplot <- ggplot(flights_weather_day, aes(x = pressure, y = arr_delay)) +
  geom_point()
lineplot <- ggplot(flights_weather_day, aes(x = pressure, y = arr_delay)) +
  geom_smooth()
plot_grid(scatterplot, lineplot, labels = "AUTO")
```

Geometric Objects 1



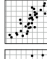
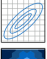


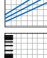

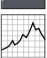






Geoms - Reference

ggplot2 provides over 40 geoms, and extensions provide more (see <https://ggplot2.tidyverse.org/reference>)

The best way to get a comprehensive overview of all the geoms is the [ggplot2 geom reference](https://posit.co/resources/ggplot2-geom-reference/), which you can find at <https://posit.co/resources/ggplot2-geom-reference/>

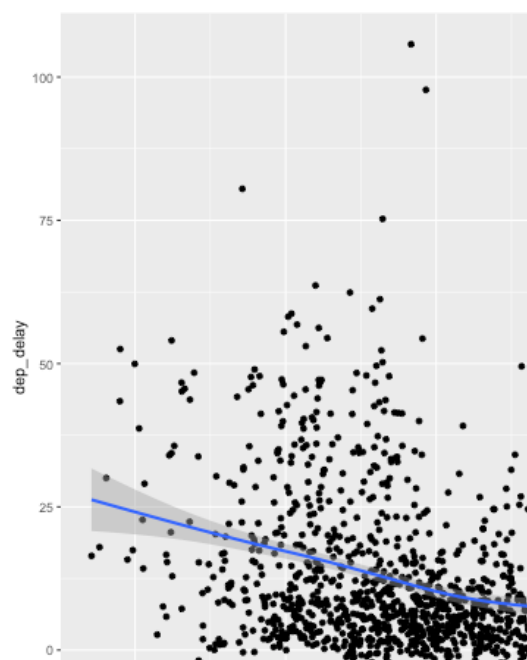
Two Variables	
Continuous X, Continuous Y <i>f</i> <- ggplot(mpg, aes(cty, hwy))	Continuous Bivariate <i>i</i> <- ggplot(movies, aes(rank, gross))
 f+ geom_blank() 	 i+ geom_bin2d() xmax, xmin, ymax, ymin, linetype, size, weight
 f+ geom_jitter() x, y, alpha, color, fill, shape, size	 i+ geom_density() x, y, alpha, colour
 f+ geom_point() x, y, alpha, color, fill, shape, size	 i+ geom_hex() x, y, alpha, colour
 f+ geom_quantile() x, y, alpha, color, linetype, size, weight	Continuous Univariate <i>j</i> <- ggplot(economics, aes(year, unemp))
 f+ geom_rug(sides = "bl") alpha, color, linetype, size	 j+ geom_area() x, y, alpha, color, fill
 f+ geom_smooth(model = lm) x, y, alpha, color, fill, linetype, size, weight	 j+ geom_line() x, y, alpha, color, fill

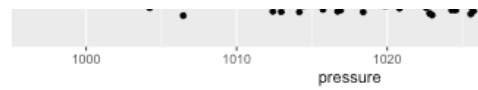
Multiple Geoms 1

To display multiple geoms in the same plot, a
`ggplot()`:

```
ggplot(flights_weather_day, aes(x = pressure, y = dep_delay)) +  
  geom_point() +  
  geom_smooth()
```

Multiple Geoms 1



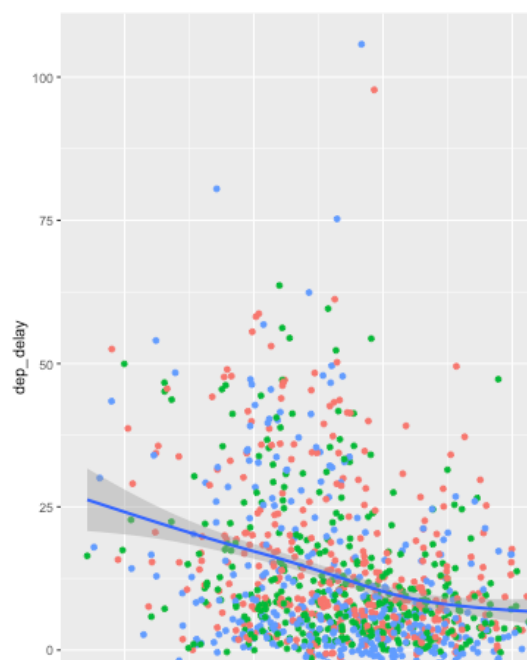


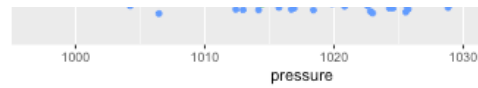
Multiple Geoms 2

If you place mappings in a `geom` function, `ggplot` will extend or overwrite the global mappings for that layer. It is possible to display different aesthetics in different layers.

```
ggplot(flights_weather_day, aes(x = pressure, y = dep_delay)) +  
  geom_point(aes(color = origin)) +  
  geom_smooth()
```

Multiple Geoms 2



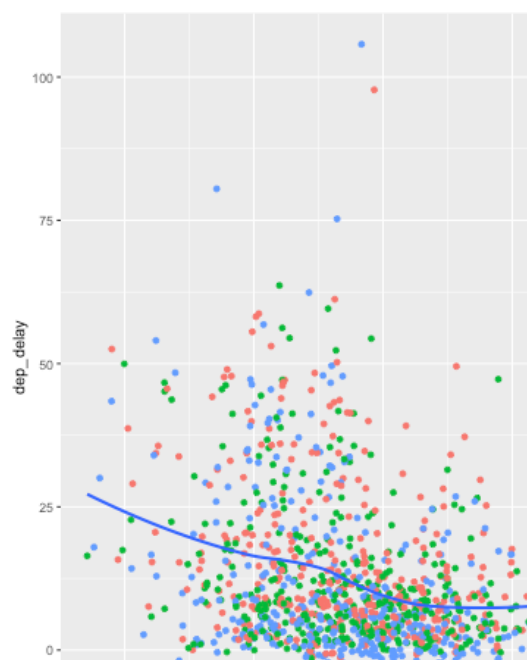


Multiple Geoms 3

You can use the same idea to specify different smooth line displays just a subset of the data JFK. The local data argument in `geom_smooth()` argument in `ggplot()` for that layer only.

```
ggplot(flights_weather_day, aes(x = pressure, y = dep_delay)) +  
  geom_point(mapping = aes(color = origin)) +  
  geom_smooth(data = filter(flights_weather_day, origin = "JFK"))
```


Multiple Geoms 3

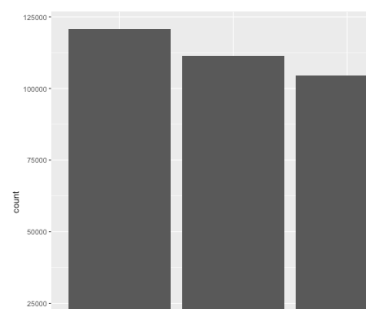


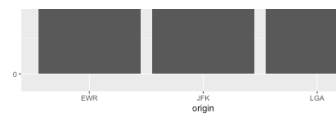


Statistical Transformations –

Let's make a bar chart of the number of flights. The algorithm uses a built-in statistical transformation to calculate the counts.

```
ggplot(flights_weather, aes(x = origin)) +  
  geom_bar()
```

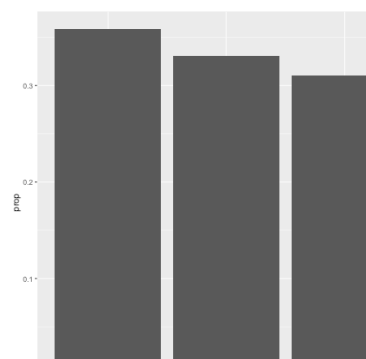




Bar charts 2

You can override the statistic a geom uses to compute the y values. If you want to plot proportions, rather than counts:

```
ggplot(flights_weather, aes(x = origin, y = stat(prop), geom_bar())
```

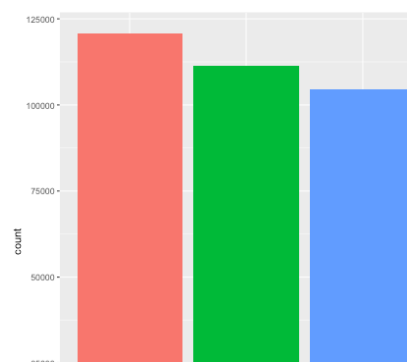


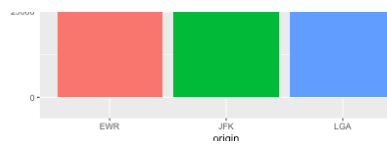


Coloring barcharts

You can color a bar chart using either the `color` (outline), or, more usefully, `fill`:

```
ggplot(flights_weather, aes(x = origin, fill= origin)) +  
  geom_bar()
```





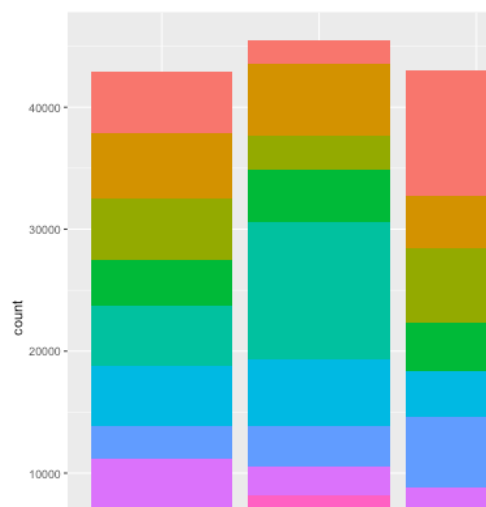
Coloring barcharts

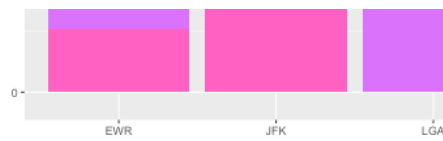
More interestingly, you can fill by another variable (e.g. to look at the destination airports with a lot of flights)

```
flights_weather_ss <- flights_weather |>
  group_by(dest) |>
  filter(n() > 10000)
```

Coloring barcharts

```
ggplot(flights_weather_ss, aes(x = origin, fill= dest)) +  
  geom_bar()
```

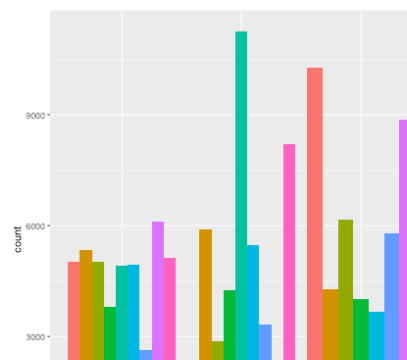


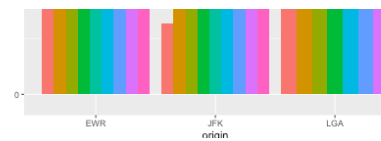


Coloring barcharts

The `position = "dodge"` places overlapping objects next to each other. This makes it easier to compare individual categories.

```
ggplot(flights_weather_ss, aes(x = origin, fill= dest)) +  
  geom_bar(position="dodge")
```



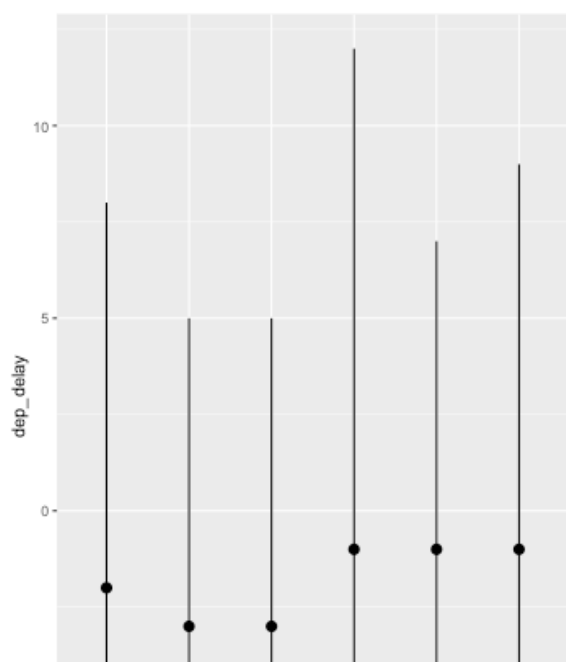


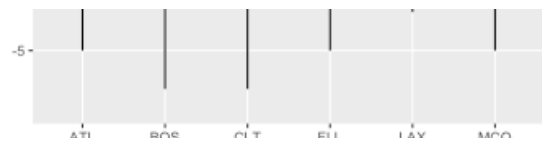
Statistical transformations -

You might want to draw greater attention to the median in your code. For example, you might use `stat_summary` to draw the y values for each unique x value, to draw attention to what you're computing:

```
ggplot(flights_weather_ss, aes(x = dest, y = dep_delay)) +  
  stat_summary(fun = median,  
              fun.min = function(z) { quantile(z,0.25) },  
              fun.max = function(z) { quantile(z,0.75) })
```


Statistical transformations -



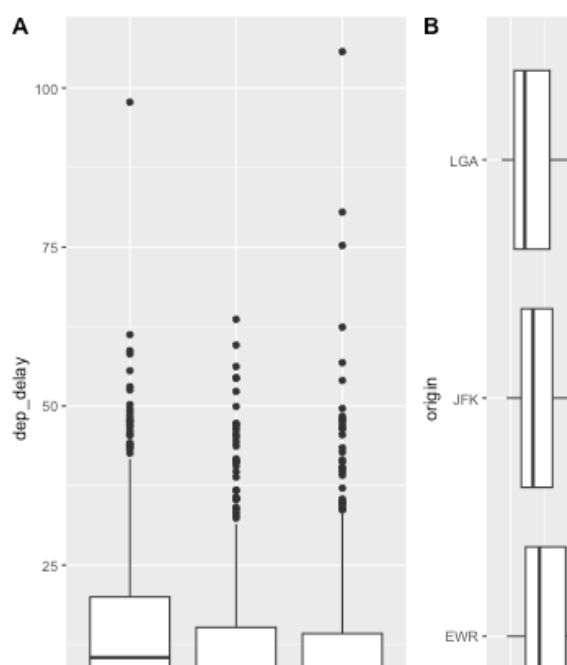


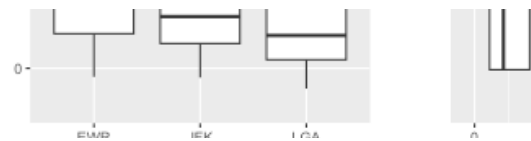
Coordinate systems

Coordinate systems are one of the more common things you can do in ggplot2. To start with something simple, here's how to flip

```
unflipped <- ggplot(flights_weather_day, aes(x = origin,
  geom_boxplot()
flipped <- ggplot(flights_weather_day, aes(x = origin, y
  geom_boxplot() +
  coord_flip()
plot_grid(unflipped, flipped, labels = "AUTO")
```

Coordinate systems





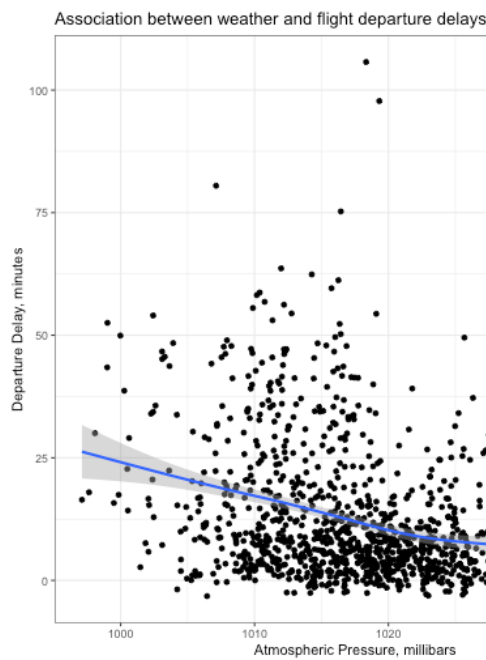
Adding labels

You can make nicer axes and add titles with `labs()`

Also showing a minimal theme that removes background grid lines with `theme_bw()`

```
ggplot(flights_weather_day, aes(x = pressure, y = dep_delay)) +  
  geom_point() +  
  geom_smooth() +  
  labs(  
    x = "Atmospheric Pressure, millibars",  
    y = "Departure Delay, minutes",  
    title = "Association between weather and flight departure delay"  
  ) +  
  theme_bw()
```

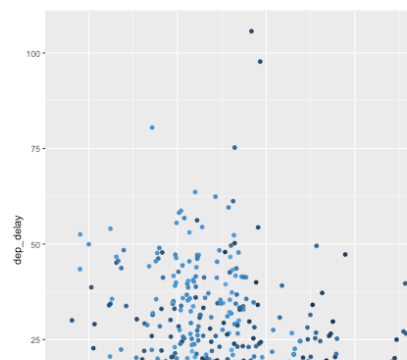
Adding labels

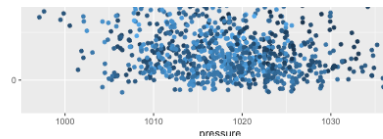


Color ramps

If you add a continuous variable in your color ramp

```
ggplot(flights_weather_day, aes(x = pressure, y = dep_delay))  
  geom_point(aes(colour = temp))
```

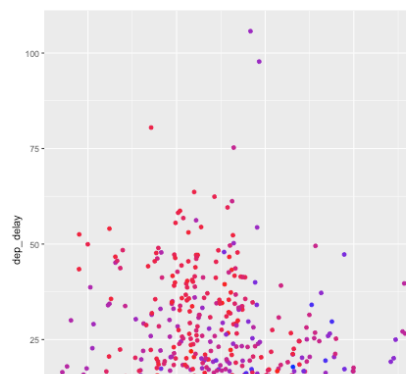


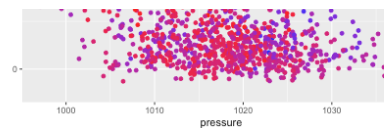


Color palettes

You can define your own color ramp or use one of the built-in ones.

```
ggplot(flights_weather_day, aes(x = pressure, y = dep_delay)) +  
  geom_point(aes(colour = temp)) +  
  scale_colour_gradient(low = "blue", high = "red")
```

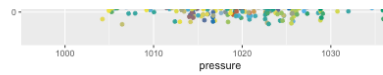




Color palettes

```
cbPalette <- c("#999999", "#E69F00", "#56B4E9", "#009E73")  
ggplot(flights_weather_day, aes(x = pressure, y = dep_delay)) +  
  geom_point(aes(colour = temp)) +  
  scale_colour_gradientn(colours = cbPalette)
```





A Great reference

A great (comprehensive) reference for everything
the R Graphics Cookbook:

<https://r-graphics.org/>

Finally, file under "useless but useful"

ggpattern - is a library for adding pattern fills

```
library(ggpattern)
df <- data.frame(level = c("a", "b", "c", "d"), outcome = c("a", "b", "c", "d"))

ggplot(df, aes(level, outcome, pattern_fill = level)) +
  geom_col_pattern(pattern = "stripe", fill = "white", color = "black") +
  theme(legend.position = "none") +
  labs(title = "ggpattern::geom_pattern_col()", subtitle = "Pattern fill by level") +
  coord_fixed(ratio = 1 / 2) +
  theme_bw() +
```

Finally, file under "useless b

