

# **Data Wrangling and M Regression**

**JSC 370: Data So**

**February 5, 202**

# Today's goals

We will learn how to wrangle and manipulate large data with `dtplyr` - i

- Selecting variables.
- Filtering data.
- Creating variables.
- Summarize data.


Throughout the session we will see examples using:

- [data.table](#) in R,
- [dtplyr](#) in R, and
- [pydatatable](#)

All with the [MET](#) dataset.

We will also take a look at advanced regression, for which you will need

# Data wrangling in R



Data wrangling describes the processes designed to import, clean up, and transform complex "raw" forms into high-quality data. You can use your wrangled

# Data wrangling in R

Overall, you will find the following approaches:

- **base R**: Use only base R functions.
- **dplyr**: Using "verbs".
- **data.table**: High-performing (ideal for large data)
- **dplyr + data.table = dtplyr**: High-performing + dplyr verbs.

Other methods involve, for example, using external tools such as [Spark](#),

We will be focusing on data.table because of [this](#)

Take a look at this very neat cheat sheet by Erik Petrovski [here](#).

## Selecting variables: Load the p

```
library(data.table)
library(dtplyr)
library(dplyr)
library(ggplot2)
library(mgcv)
library(lubridate)
```

The dtplyr R package translates dplyr (tidyverse) syntax to data table **verbs** while at the same time leveraging the performance of data.table.

The mgcv package enables advanced regression models with basis splines.

## Loading the data

We will use the MET dataset, which we can download (and load) directly.

```
# Where are we getting the data from
met_url <- "https://raw.githubusercontent.com/JSC370/JSC370/master/data/MET/MET.gz"

# Downloading the data to a tempfile (so it is destroyed
# you can replace this with, for example, your own data:
# tmp <- tempfile(fileext = ".gz")
tmp <- "met.gz"

# We could be downloading this, ONLY IF this was not downloaded
# otherwise is just a waste of time.
if (!file.exists(tmp)) {
  download.file(
    url      = met_url,
    destfile = tmp,
    # method  = "libcurl", timeout = 1000 (you may need to increase this)
  )
}
```

Now we can load the data using the `fread()` function.

## Reading in the data

In R, `fread`, do a quick wrangle to remove outliers (discovered earlier), and

```
met_dt <- fread(tmp)
met_dt <- met_dt[temp > -10][order(temp)]
head(met_dt)
```

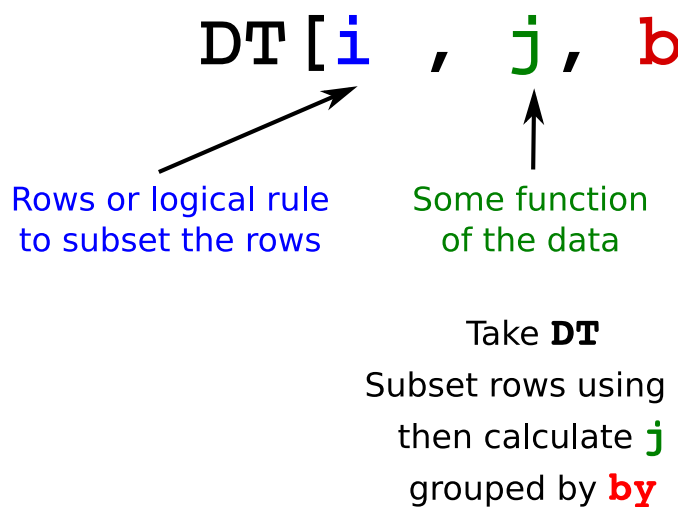
In Python, import with `datatable`, read, and print first 5 rows

```
import datatable as dt
met_dt_py = dt.fread("met.gz")
met_dt_py.head(5)
```

Before we continue, let's learn a bit more on `data.table` and `dplyr`

## data.table and dtplyr: Data

- As you have seen in previous lectures, in data.table all happens with  
imagine DT:



- Any time that you see **:=** in **j** that is "Assignment by reference." Using =



## data.table and dtplyr: Data

Operations applied in `j` are evaluated *within* the data, meaning that names


```
# This returns an error (met is not referencing the data.)  
met[, elev]
```

```
# This works fine  
met_dt[, elev]
```

# data.table and dtplyr: Data

Furthermore, we can do things like this:

```
met_dt[, plot(temp, elev)]
```



## NULL

## Lazy loading, queries

- From [Wikipedia](#) "Lazy Loading" (also known as asynchronous loading) is a technique in computer programming and mostly in web design and development where data is loaded at which it is needed. It can contribute to efficiency in the program's execution.
- Lazy loading means that the code for a particular function doesn't run until the last minute – when it's actually being used.
- When you create a "lazy" query, you're creating a pointer to a set of data that won't actually run and the data isn't actually loaded until you call "next" or "load" and load it into an object.

## **data.table and dtplyr: Lazy**

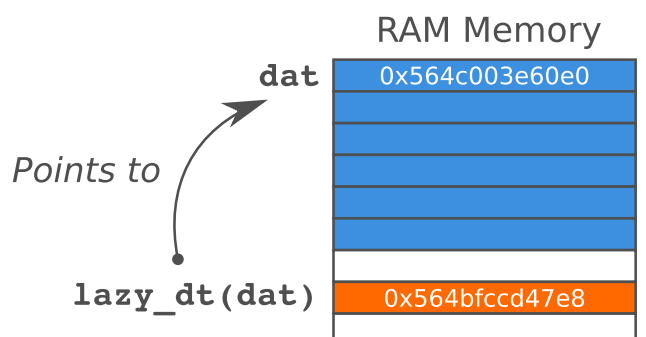
- The dtplyr package provides a way to translate dplyr verbs to data.table
- The key lies on the function lazy\_dt from dtplyr (see ?dtplyr)
- This function creates a wrapper that "points" to a data.table object

## data.table and dtplyr: Lazy

```
# Creating a lazy table object
met_ldt <- lazy_dt(met_dt, immutable = FALSE)

# We can use the address() function from data.table
address(met_ldt)
address(met_ldt$parent)

## [1] "0x7fcbff8c9788"
## [1] "0x7fcbe5255c00"
```




# data.table selecting columns

How can we select the columns USAFID, lat, and lon, using data.table names:

```
met_dt[, list(USAFID, lat, lon, temp, elev)]
# met_dt[, .(USAFID, lat, lon, temp, elev)] # Alternative
# met_dt[, c("USAFID", "lat", "lon", "temp", "elev")] # Another
```

	USAFID	lat	lon	temp	elev
##	1: 726764	44.683	-111.116	-3.0	2025
##	2: 726764	44.683	-111.116	-3.0	2025
##	3: 726764	44.683	-111.116	-3.0	2025
##	4: 726764	44.683	-111.116	-3.0	2025
##	5: 720411	36.422	-105.290	-2.4	2554
##	---				
##	2317200: 690150	34.300	-116.166	52.8	696
##	2317201: 690150	34.296	-116.162	52.8	625
##	2317202: 690150	34.300	-116.166	53.9	696
##	2317203: 690150	34.300	-116.166	54.4	696
##	2317204: 720267	38.955	-121.081	56.0	467

What happens if instead of `list()` you used `c()`?

## Selecting columns (cont. 1)

Using the **dplyr::select** verb:

```
met_dt |>
  select(USAFID, lat, lon, temp, elev)
```

```
##           USAFID      lat      lon temp elev
##      1: 726764 44.683 -111.116 -3.0 2025
##      2: 726764 44.683 -111.116 -3.0 2025
##      3: 726764 44.683 -111.116 -3.0 2025
##      4: 726764 44.683 -111.116 -3.0 2025
##      5: 720411 36.422 -105.290 -2.4 2554
##      ---
## 2317200: 690150 34.300 -116.166 52.8  696
## 2317201: 690150 34.296 -116.162 52.8  625
## 2317202: 690150 34.300 -116.166 53.9  696
## 2317203: 690150 34.300 -116.166 54.4  696
## 2317204: 720267 38.955 -121.081 56.0  467
```

## Selecting columns (cont. 2)

In the case of pydatatable

```
met_dt_py[:, ["USAFID", "lat", "lon", "temp", "elev"]]
```

What happens if instead of ["USAFID", "lat", "lon", "temp", "lon", "temp", "elev"]} (vector vs set).



## Selecting columns (cont. 3)

For the rest of the session we will be using these variables: USAFID, WBAN, wind.sp, temp, and atm.press.

```
# Data.table
met_dt <- met_dt[,
  .(USAFID, WBAN, year, month, day,
    hour, min, lat, lon, elev,
    wind.sp, temp, atm.press)
]

# Need to redo the lazy table
met_ldt <- lazy_dt(met_dt)
```

## Data filtering: Logical condition

- Based on logical operations, e.g. condition 1 [and|or condition]
- Need to be aware of ordering and grouping of and and or operators
- Fundamental **logical** operators:

x	y	Negate !x	And x & y	Or x   y	Xor xor(x, y)
true	true	false	true	true	false
false	true	true	false	true	true
true	false	false	false	true	true
false	false	true	false	false	false

- Fundamental **relational** operators, in R: <, >, <=, >=, ==, !=.

# XOR operations

- The [XOR logical operation](#), exclusive or, takes two Boolean operands and returns true if the two operands are different. Conversely, it returns false if the two operands have the same value.
- So, for example, the XOR operator can be used when we have to check if two conditions are true at the same time.

# How many ways can you write

Write a function that takes two arguments ( $x, y$ ) and applies the XOR operation. Here is a template:

```
myxor <- function(x, y) {  
  res <- logical(length(x))  
  for (i in 1:length(x)) {  
    res[i] <- # do something with x[i] and y[i]  
  }  
  return(res)  
}
```

Or if vectorized (this would be better)

```
myxor <- function(x, y) {  
  # INSERT YOUR CODE HERE  
}
```

Hint 1: Remember that negating ( $x \& y$ ) equals  $(!x \mid !y)$ .

Hint 2: Logical operators are distributive, meaning  $a * (b + c) = (a * b) + (a * c)$ .

In R

```
myxor1 <- function(x,y) {(x & !y) | (!x & y)}
myxor2 <- function(x,y) {!( (!x | y) & (x | !y))}
myxor3 <- function(x,y) {(x | y) & (!x | !y)}
myxor4 <- function(x,y) {!( (!x & !y) | (x & y))}
cbind(
  ifelse(xor(test[,1], test[,2]), "true", "false"),
  ifelse(myxor1(test[,1], test[,2]), "true", "false"),
  ifelse(myxor2(test[,1], test[,2]), "true", "false"),
  ifelse(myxor3(test[,1], test[,2]), "true", "false"),
  ifelse(myxor4(test[,1], test[,2]), "true", "false")
)

##      [,1]      [,2]      [,3]      [,4]      [,5]
## [1,] "false" "false" "false" "false" "false"
## [2,] "true"  "true"  "true"  "true"  "true"
## [3,] "true"  "true"  "true"  "true"  "true"
## [4,] "false" "false" "false" "false" "false"
```

Or in Python

```
# Loading the libraries
import numpy as np
import pandas as pa

# Defining the data
x = [True, True, False, False]
y = [False, True, True, False]
ans = {
    'x' : x,
    'y' : y,
    'and' : np.logical_and(x, y),
    'or' : np.logical_or(x, y),
    'xor' : np.logical_xor(x, y)
}
pa.DataFrame(ans)
```

Or in Python (bis)

```
def myxor(x,y):  
    return np.logical_or(  
        np.logical_and(x, np.logical_not(y)),  
        np.logical_and(np.logical_not(x), y)  
    )  
  
ans['myxor'] = myxor(x,y)  
pa.DataFrame(ans)
```

We will now see applications using the met dataset

# Filtering (subsetting) the data

Need to select records according to some criteria. For example:

- First day of the month, and
- Above latitude 40, and
- Elevation outside the range 500 and 1,000.

The logical expressions would be

- `(day == 1)`
- `(lat > 40)`
- `((elev < 500) | (elev > 1000))`

Respectively.



In R with `data.table`:

```
met_dt[(day == 1) & (lat > 40) & ((elev < 500) | (elev >
  nrow()))]
```

```
## [1] 27049
```

In R with **dplyr::filter()**:

```
met_ldt |>
  filter(day == 1, lat > 40, (elev < 500) | (elev > 1000))
  collect() |> # Notice this line!
  nrow()
```

```
## [1] 27049
```

With lazy tables, R delays doing any work until the last possible moment and then sends it to the database in one step.

In Python

```
met_dt_py[(dt.f.day == 1) & (dt.f.lat > 40) & ((dt.f.elev  
# met_dt_py[dt.f.day == 1,:][dt.f.lat > 40,:][dt.f.elev
```

In the case of pydatatable we use `dt.f.` to refer to a column. `df.` is wh

The [f. is a symbol](#) that allows accessing column names in a datatable's

## More wrangling questions

1. How many records have a temperature within 18 and 25 C?
2. Some records have missings. Count how many records have temp
3. Following the previous question, plot a sample of 1,000 of ( lat, ' with data.

# Solutions

```
# Question 1
message("Question 1: ", nrow(met_dt[(temp < 25) & (temp > 25)]))

## Question 1: 908047

# met_dt[temp %between% c(18, 25), .N]

# met_ldt |>
#   filter(between(temp, 18, 25)) |>
#   collect() |>
#   nrow()

# Question 2
message("Question 2: ", met_dt[is.na(temp), .N])

## Question 2: 60089
```

- Note the special symbol `.N` in `j`

- `.N` can be used in `j`, which is particularly useful to get the number


## Solutions (con't)

```
# Question 3
set.seed(123)
message("Question 3")

# Drawing a sample
idx <- met_dt[, list(x = sample.int(.N, 2000, replace = F

# Visualizing the data
ggplot(map_data("state"), aes(x = long, y = lat)) +
  geom_map(aes(map_id = region), map = map_data("state"),
  geom_jitter(
    data      = met_dt[idx],
    mapping = aes(x = lon, y = lat, col = is.na(temp)),
    inherit.aes = FALSE, alpha = .5, cex = 2
  )
```

## Solutions (con't)



# Creating variables: Data types

- **logical:** Bool true/false type, e.g. dead/alive, sick/healthy, good/bad
- **strings:** string of characters (letters/symbols), e.g. names, text, etc
- **integer:** Numeric variable with no decimal (discrete), e.g. age, days
- **double:** Numeric variable with decimals (continuous), e.g. distance

In C (and other languages), strings, integers, and doubles may be specified as 8, 16, and 32 bits. This is relevant when managing large datasets, where



# Creating variables: Special data

Most programming languages have special types which are built using built-in functions.

- **time:** Could be date, date + time, or a combination of both. Usually, in R, the `Date` class has as reference 1970-01-01, in other words, "date of epoch".
- **categorical:** Commonly used to represent strata/levels of variables. It is stored as a factor, where the data is stored as numbers but has a label.
- **ordinal:** Similar to factor, but it has ordering, e.g. "satisfaction level".

Other special data types could be ways to represent missings (usually denoted by `NA`), e.g. `+-Inf` and `Undefined (NaN)`.

When storing/sharing datasets, it is a good practice to do it along a dictionary.

## Questions 3: What's the best v following

- 0, 1, 1, 0, 0, 1
- Diabetes type 1, Diabetes type 2, Diabetes type 1, Diabetes type 2
- on, off, off, on, on, on
- 5, 10, 1, 15, 0, 0, 1
- 1.0, 2.0, 10.0, 6.0
- high, low, medium, medium, high
- -1, 1, -1, -1, 1,
- .2, 1.5, .8,  $\pi$
- $\pi$ ,  $\exp 1$ ,  $\pi$ ,  $\pi$

# Variable creation

If we wanted to create two variables,  $elev^2$  and the scaled version of  $wind.sp$ , we can do the following

With `data.table`

```
met_dt[, elev2 := elev^2]
met_dt[, windsp_scaled := wind.sp/sd(wind.sp, na.rm = TRUE)]

# Alternatively:
# met_dt[, c("elev2", "windsp_scaled") := .(elev^2, wind.sp/sd(wind.sp, na.rm = TRUE))]
```

## Variable creation (cont. 1)

With the verb **dplyr::mutate()**:

```
met_dt[, c("elev2", "windsp_scaled") := NULL] # This to d
met_ldt |>
  mutate(
    elev2 = elev ^ 2,
    windsp_scaled = wind.sp/sd(wind.sp,na.rm=TRUE)
  ) |>
  collect()
```

```
## # A tibble: 2,317,204 × 15
##   USAFID WBAN  year month  day  hour  min  lat  lon
##   <int> <int> <int> <int> <int> <int> <int> <dbl> <dbl>
## 1 726764 94163 2019     8    27    11    50  44.7 -111
## 2 726764 94163 2019     8    27    12    10  44.7 -111
## 3 726764 94163 2019     8    27    12    30  44.7 -111
## 4 726764 94163 2019     8    27    12    50  44.7 -111
## 5 720411   137 2019     8    18    12    35  36.4 -108
## 6 726764 94163 2019     8    26    12    30  44.7 -111
```

```
## 7 726764 94163 2019      8      26      12      50  44.7 -11
## 8 726764 94163 2019      8      26      13      10  44.7 -11
## 9 726764 94163 2019      8      27      10      30  44.7 -11
## 10 726764 94163 2019     8      27      10      50  44.7 -11
```

## Variable creation (cont. 2)

Imagine that we needed to generate all those calculations (scale by sd)  
**.SD** symbol:

```
# Listing the names
in_names <- c("wind.sp", "temp", "atm.press")
out_names <- paste0(in_names, "_scaled")
met_dt[,
  c(out_names) := lapply(.SD, function(x) x/sd(x, na.rm = TRUE))
  .SDcols = in_names
]

# Looking at the first 4
head(met_dt[, .SD, .SDcols = out_names], n = 4)
```

```
##      wind.sp_scaled temp_scaled atm.press_scaled
## 1:              0 -0.4955951              NA
## 2:              0 -0.4955951              NA
## 3:              0 -0.4955951              NA
## 4:              0 -0.4955951              NA
```

- Key things to notice here: **c(out\_names)**, **.SD**, and **.SDCols**.
- More on [.SD](#)

## Variable creation (cont. 3)

In the case of dplyr, we could use the following

```
as_tibble(met_ldt) |>
  mutate(
    across(
      all_of(in_names),
      function(x) x/sd(x, na.rm = TRUE),
      .names = "{col}_scaled2"
    )
  ) |>
  # Just to print the last columns
  select(ends_with("_scaled2")) |>
  head(n = 4)
```

```
## # A tibble: 4 × 3
##   wind.sp_scaled2 temp_scaled2 atm.press_scaled2
##             <dbl>         <dbl>          <dbl>
## 1              0         -0.496             NA
## 2              0         -0.496             NA
```

##	3	0	-0.496	NA
##	4	0	-0.496	NA

Key thing here: This approach has no direct translation to `data.table`

## Merging data

- While building the MET dataset, we dropped the State data.
- We can use the original Stations dataset and *merge* it to the MET d
- But we cannot do it right away. We need to process the data somev

## Merging data (cont. 1)

```
stations <- fread("ftp://ftp.ncdc.noaa.gov/pub/data/noaa/
stations[, USAF := as.integer(USAF)]

# Dealing with NAs and 999999
stations[, USAF := fifelse(USAF == 999999, NA_integer_,
stations[, CTRY := fifelse(CTRY == "", NA_character_, C
stations[, STATE := fifelse(STATE == "", NA_character_,

# Selecting the three relevant columns, and keeping unique
stations <- unique(stations[, list(USAF, CTRY, STATE)])

# Dropping NAs
stations <- stations[!is.na(USAF)]

head(stations, n = 4)

##      USAF CTRY STATE
## 1: 7018 <NA>  <NA>
## 2: 7026  AF  <NA>
## 3: 7070  AF  <NA>
```



```
## 4: 8260 <NA> <NA>
```

Notice the function `fifelse()`. Now, let's try to merge the data!

## Merging data (cont. 2)

```
merge(  
  # Data  
  x      = met_dt,  
  y      = stations,  
  # List of variables to match  
  by.x   = "USAFID",  
  by.y   = "USAF",  
  # Which obs to keep?  
  all.x  = TRUE,  
  all.y  = FALSE  
) |> nrow()
```

```
## [1] 2385443
```

This is more rows! The original dataset, `met_dt`, has 2317204. This means we have duplicate IDs. We can fix this:

```
stations[, n := 1:.N, by = .(USAF)]
stations <- stations[n == 1,][, n := NULL]
```

## Merging data (cont. 3)

We now can use the function `merge()` to add the extra data

```
met_dt <- merge(
  # Data
  x = met_dt,
  y = stations,
  # List of variables to match
  by.x = "USAFID",
  by.y = "USAF",
  # Which obs to keep?
  all.x = TRUE,
  all.y = FALSE
)

head(met_dt[, list(USAFID, WBAN, STATE)], n = 4)

##      USAFID  WBAN STATE
## 1: 690150 93121    CA
## 2: 690150 93121    CA
```

```
## 3: 690150 93121 CA
## 4: 690150 93121 CA
```

What happens when you change the options `all.x` and `all.y`?

## Aggregating data: Adding group

- Many times we need to either impute some data, or generate variables.
- If we, for example, wanted to impute missing temperature with the mean of the temperature by state, we can do this with the **`data.table::fcoalesce()`** function:

```
met_dt[, temp_imp := fcoalesce(temp, mean(temp, na.rm = TRUE),
  by = .(STATE, year, month, day))]
```

- In the case of dplyr, we can do the following using **`dplyr::group_by`** together with **`mutate()`**:

```
# We need to create the lazy table again, since we replaced the data
met_ldt <- lazy_dt(met_dt, immutable = FALSE)

met_ldt |>
  group_by(STATE, year, month, day) |>
  mutate(
    temp_imp2 = coalesce(temp, mean(temp, na.rm = TRUE),
      by = .(STATE, year, month, day))
  ) |> collect()
```

## Aggregating data: Adding group (cont.)


Let's see how it looks like

```
# Preparing for ggplot2
plotdata <- met_dt[USAFID == 720172][order(year, month, day)]
plotdata <- rbind(
  plotdata[, .(temp = temp, type = "raw")],
  plotdata[USAFID == 720172][, .(temp = temp_imp, type = "imp")]
)

# Generating an 'x' variable for time
plotdata[, id := 1:.N, by = type]

plotdata |>
  ggplot(aes(x = id, y = temp, col = type, lty = type)) +
  geom_line()
```

## Aggregating data: Adding group (cont.)





## Aggregating data: Summary ta

- Using `by` also allow us creating summaries of our data.
- For example, if we wanted to compute the average temperature, wind speed, and atmospheric pressure, we could do the following

```
met_dt[, .(  
  temp_avg      = mean(temp, na.rm=TRUE),  
  wind.sp_avg   = mean(wind.sp, na.rm=TRUE),  
  atm.press_avg = mean(atm.press, na.rm = TRUE)  
)  
  by = STATE  
][order(STATE)] |> head(n = 4)
```

```
##      STATE temp_avg wind.sp_avg atm.press_avg  
## 1:    AL 26.19799    1.563645    1016.148  
## 2:    AR 26.20697    1.872876    1014.551  
## 3:    AZ 28.80596    2.983999    1010.771  
## 4:    CA 22.36199    2.614711    1012.637
```

## Aggregating data: Summary ta

When dealing with too many variables, we can use the `.SD` special sym

```
# Listing the names
in_names <- c("wind.sp", "temp", "atm.press")
out_names <- paste0(in_names, "_avg")

met_dt[,
  setNames(lapply(.SD, mean, na.rm = TRUE), out_names),
  .SDcols = in_names, keyby = STATE
] |> head(n = 4)
```

##	STATE	wind.sp_avg	temp_avg	atm.press_avg
## 1:	AL	1.563645	26.19799	1016.148
## 2:	AR	1.872876	26.20697	1014.551
## 3:	AZ	2.983999	28.80596	1010.771
## 4:	CA	2.614711	22.36199	1012.637

Notice the **keyby** option here: "Group by STATE and order by STATE".

## Aggregating data: Summary ta

- Using **dplyr** verbs

```
met_ldt |>
  group_by(STATE) |>
  summarise(
    temp_avg      = mean(temp, na.rm=TRUE),
    wind.sp_avg   = mean(wind.sp, na.rm=TRUE),
    atm.press_avg = mean(atm.press, na.rm = TRUE)
  ) |> arrange(STATE) |> head(n = 4)

## Source: local data table [4 x 4]
## Call:   head(setorder(`_DT3`[, .(temp_avg = mean(temp,
##     wind.sp_avg = mean(wind.sp, na.rm = TRUE), atm.press
##     na.rm = TRUE)), keyby = .(STATE)], STATE, na.la
##     n = 4)
##
##   STATE temp_avg wind.sp_avg atm.press_avg
##   <chr>   <dbl>      <dbl>        <dbl>
## 1 AL      26.2        1.56         1016.
```



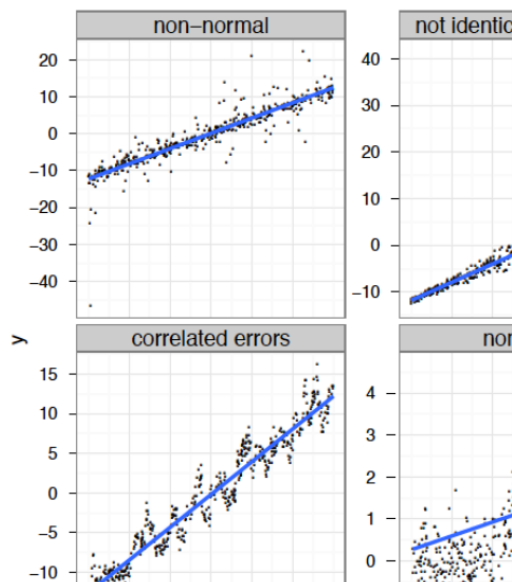
```
## 2 AR      26.2      1.87      1015.
## 3 AZ      28.8      2.98      1011.
## 4 CA      22.4      2.61      1013.
##
```


## Other data.table goodies

- `shift()` Fast lead/lag for vectors and lists.
- `fifelse()` Fast if-else, similar to base R's `ifelse()`.
- `fcoalesce()` Fast coalescing of missing values.
- `%between%` A short form of `(x < lb) & (x > up)`
- `%inrange%` A short form of `x %in% lb:up`
- `%chin%` Fast match of character vectors, equivalent to `x %in% X`,
- `nafill()` Fill missing values using a constant, last observed value

# Machine Learning 1: Advanced

- Linear regression is useful, but there are so many ways in which it can fail





# Machine Learning 1: Advanced

- A linear model tries to fit the best straight line that passes through
- In general,  $Y(s) = f(s) + \epsilon$  where in regular linear regression  $f(s)$
- If we want to represent the regression more generally, we can define a basis function consisting of 'non-linear' terms.

# Basis Function

## Basics of Basis Functions

- We will start with a 1-dimensional, univariate case. For example this could be time (x) with basis functions.
- Polynomial bases are a good way to illustrate what is going on. Consider

$$y_i = f(x_i) + \epsilon_i$$

and let's expand it out by a polynomial

$$y_i = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \beta_3 x_i^3 + \dots$$

# Basis Function

Here

$$f(x_i) = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \beta_3$$

is a 4th order polynomial. So,  $f(x)$  is a function represented by **five** bas

$$f(x_i) = \sum_{j=1}^5 x^j \beta_j = \sum_{j=1}^5 b_j$$

that are defined by:

$$b_1(x) = 1, b_2(x) = x, b_3(x) = x^2, b_4(x)$$

# Basis Functions

- In general, a basis is a set of functions that can be added together to form a complicated function
- Here our weights are the regression coefficients  $\beta_j$
- In general, a basis function is represented by

$$f_i = \sum b_j(x_i)\beta_j$$

$$\begin{pmatrix} f_1 \\ f_2 \\ f_3 \\ f_4 \\ f_5 \end{pmatrix} = \begin{bmatrix} 1 & b_1(x_1) & b_2(x_1) & b_3(x_1) & b_4(x_1) & b_5(x_1) \\ 1 & b_1(x_2) & b_2(x_2) & b_3(x_2) & b_4(x_2) & b_5(x_2) \\ 1 & b_1(x_3) & b_2(x_3) & b_3(x_3) & b_4(x_3) & b_5(x_3) \\ 1 & b_1(x_4) & b_2(x_4) & b_3(x_4) & b_4(x_4) & b_5(x_4) \\ 1 & b_1(x_5) & b_2(x_5) & b_3(x_5) & b_4(x_5) & b_5(x_5) \end{bmatrix}$$

# Polynomial Basis





Figure 3.1 Illustration of the idea of representing a function in a polynomial basis. The first 5 panels (starting from top left),  $b_i(x)$  for a 4th order polynomial basis. The basis functions


$\phi_j(x)$ , for a 4th order polynomial basis. The basis functions are multiplied by the parameter,  $\beta_j$ , and are then summed to give the final curve. The curve is shown in the bottom right panel. By varying the  $\beta_j$ , we can vary the curve to fit any polynomial function of order 4 or lower. See also figure 3.

## Polynomial Basis

- The basis functions are each multiplied by  $\beta_j$  and then summed to give the final curve, as shown in the bottom left figure.
- Below, we show this concept in terms of an example of CO<sub>2</sub> concentration.








# Splines


- In general, splines are curves that are formed by combining pieces
- There are several types of splines including natural, cubic, and b-splines

$$f(t_i) = \sum_{j=1}^4 t^j \beta_j$$

- B-spline curves are made up of polynomial pieces and are defined
- Choosing the number of knots defines how smooth (few) or wiggly

# Splines





# Splines

- Smoothing splines with penalty allows us to estimate where to put the knots
- Minimize the function

$$\sum_i (y_i - f(t_i))^2 + \lambda \int f''^2$$

- Here,  $\lambda$  is a penalty parameter that controls how much to penalize wiggly
- Trade-off between the goodness of fit (the sum of squares) and the wiggli
- Start by putting a knot at every data point, then penalize.
- It is an optimization problem where we minimize:

$$\sum_i (y_i - B_i^T \beta)^2 + \lambda \beta^T S \beta$$

- the matrix S is constructed by using the spline basis we chose, B is the ba

# Splines

- This function

$$\sum_i (y_i - f(t_i))^2 + \lambda \int f''(t)^2 dt$$

represents the loss + penalty.

- We see similar functions in lasso and ridge regression.
- The second derivative  $f''(t)^2$  corresponds to how much the slope is changing (the slope of the function at  $t$ ).
- The integral  $\int f''(t)^2 dt$  is a measure of the total change in the function  $f'$ . If  $f'(t)$  will be close to constant and the integral will take on a small value.
- A large  $\lambda$  will make the function  $f$  smoother, but  $\lambda = 0$  means the function is as wiggly as possible.
- As  $\lambda \rightarrow \infty$ ,  $f$  will be perfectly smooth, a straight line that passes through the mean of the data.

# 1-D Splines

Types of 1-D splines include:

- cubic splines (basically piecewise cubic polynomials)
- cyclic splines (cubic splines with connected ends)
- basis splines (B-spline) with other polynomial orders
- cardinal splines (where knot placement is always a certain distance)
- wavelets (often cardinal wavelet splines)

# Fitting Spline Regression Model

We will use CO<sub>2</sub> data from the Mauna Loa observatory in Hawaii: [https://scrippsco2.ucsd.edu/](#)


- important variables are: average (monthly CO<sub>2</sub> concentrations)
- we will make a month-year variable

```
co2 <- read.csv("co2_mm_mlo.csv", skip=40)

co2 <- co2 |>
  mutate(month_year = make_date(co2$year, co2$month)) |>
  rename(co2 = average)


co2 |>
  ggplot(aes(y=co2,x=month_year)) +
  geom_line() +
  labs(x='Date (month-year)', y='CO2 concentration ppm')+
  theme_bw()
```

# Fitting Spline Regression Model



# Fitting Spline Regression Model

```
library(mgcv)
# Using cubic regression spline bases with 4 knots to show
co2_2023 <- co2[co2$year==2023,]
gam_co2 <- gam(co2~s(month,bs="cr", k=4),data=co2_2023)
plot(gam_co2)
```






2 4 6 8  
month

## Fitting Spline Regression Models

```
# try fitting to all data and smoothing date (overall trend)
gam_co2_all <- gam(co2~s(decimal.date,bs="cr",k=20))+s(month)
# predict on data
pred_co2 <- predict.gam(gam_co2_all,co2)
plot(pred_co2,type='l')
```



## 2-D Splines

- Thin plate splines are smoothing splines in 2-d
- Extend the 1-d case to:


$$\sum_i (z_i - g(s_1, s_2))^2 + \lambda \iint g''(s_1, s_2)$$


- where the penalty breaks down to the sum of the partial second derivatives
- $\lambda$  controls the "wiggleness" as in the 1-D spline (roughness penalty)

# Thin Plate Splines


The idea behind a thin plate spline is:

- Basically we put a bendable plane through over the space and the points in
- Where there are more points grouped, we expect the plane to be pulled more
- If there is a very bumpy surface, there will be more knots used and a more






# Thin Plate Splines




# Thin Plate Spline Regression

The height of where the surface is pulled is going to depend on the mag





# Thin Plate Spline Regression



## Fitting Spline Regression Models

```
gam_temp <- gam(temp~s(x,y,bs="ts",k=60, fx=TRUE),data=id)  
plot(gam_temp)  
summary(gam_temp)
```





## More on Advanced Regression

For more information and examples about regression that includes basis functions, see the book [Statistical Learning with applications in R](#)



